

CARL SHEN

github.com/cshen98
linkedin.com/in/cshen98

cshen98.github.io

+1 (226) 700-3906
carl.shen@uwaterloo.ca

SKILLS

Languages: Java, C++, JavaScript, Python; Distributed Computing: Hadoop, Spark; Stream Processing: Kafka, Storm; Datastores & NoSQL: SQL, Redis, Druid; ETL/Data Pipelines: Gobblin, Data Factory; Machine Learning: Weka, MLlib

EXPERIENCE

Facebook - Core Payments Infrastructure Team - Software Engineering Intern Menlo Park, CA | May - Aug 2019

- The Payments team powers payments/financial related features and services across Facebook and Instagram.

LinkedIn - ETL/Apache Gobblin Team - Software Engineering Intern Mountain View, CA | Sep - Dec 2018

- Enhanced a *MapReduce*-based *Gobblin* Distcp flow to allow for splitting of files into block level granularity work units, alleviating a mapper skew bottleneck by increasing parallelization of large files; resulted in up to *~4x* performance gain in data movement for various pipelines that supported various LinkedIn-Microsoft product integrations.

Oath/Yahoo - Daily Fantasy Sports Backend Team - Software Engineering Intern Sunnyvale, CA | Jan - Apr 2018

- Implemented a heuristic-based projected points algorithm using a *Storm* topology and leveraging a *Redis* datastore, powering a set of new live projections features within the live contest details page of the Daily Fantasy platform; resulted in increased user engagement (*~5%* more page views) during live contests.
- Built and architected *MySQL/Datanucleus JDO* schemas and *Java JAX-RS/Jersey* APIs to develop components of the backend infrastructure for a new product, Yahoo Fantasy Slate, which had *~25k* users within weeks of launching.

Communications Research Centre Canada - Software Engineering Intern Ottawa, ON | Sep - Dec 2016

- Created an automated batch processing pipeline for a city wide radio spectrum monitoring system, which involved porting and integrating *Spark* data analytics scripts to be invoked by a *Data Factory* workflow via calls to a *Livy* REST interface; reduced execution time compared to existing process by up to *~7x*.

EDUCATION

University of Waterloo - BSE Candidate - 3A Software Engineering Waterloo, ON | Expected Apr 2021

- Overall GPA - *3.98/4* | Cumulative Average - *93.04%*
- Awards - First in Class/Upper Year Faculty Scholarship (\$500 | 2018), Microsoft Tuition Scholarship (\$2000 | 2017), President's Research Award (\$1500 | 2017), Madter Engineering Faculty Entrance Scholarship (\$5500 | 2015)
- Undergraduate Research Assistantships - Distributed Systems Lab, Stochastic Decoding Group
- Extracurriculars - IEEE CSCWD Conference (2016) Organizing Team, London Science Fair Judging & Organizing Roles

PROJECTS & PUBLICATIONS

Spark Hockey Analytics 2018

- Using *Kafka* and *Spark* to process data from NHL APIs for hockey analytics related projects, including a Redzone-like streaming script for predicting goal probabilities of each currently live game.

WiFi-Based Indoor Localization System (IEEE CSCWD 2017 | DOI 19.1109/CSCWD.2017.806671Z) 2017

- Implemented a localization approach based on WiFi signal strengths info relative to mobile beacons only, reducing setup costs compared to a traditional fingerprint approach while achieving a similar localization accuracy.

User Behavior Prediction Model (Cluster Computing 2017 Vol 20 Issue 2 | DOI 10.1007/s10586-017-0749-z) 2016

- Developed components of a novel smart home user behavior prediction model using *MapReduce* which provided improved execution speeds by a factor of up to *~5x*.

Smart Bed Monitoring System 2016

- Built a system to recognize and monitor bed/sleep related activity by implementing a decision tree generated using a *Weka* machine learning library training algorithm on an experimentally gathered data set.